

Introduction to Machine Learning

Lecture 3: Regression

Alexis Zubiolo

`alexis.zubiolo@gmail.com`

Data Science Team Lead @ Adcash

November 3, 2016

Before we start

Would you be interested in a **more advanced course**? I can propose

- ▶ Machine learning **from scratch** (how to implement an ML algorithm with no library)
- ▶ A **more advanced version** of this course (with more theoretical technical details)
- ▶ **Large-scale** machine learning (distributed computing)

Regression in Machine Learning

This lecture is about **regression** in Machine learning.

Reminder: In regression, the output y is **continuous**.

Example:

- ▶ **Price estimation:** $y = \text{price}$ (e.g. 50000 BGN for a house)
- ▶ **Predicting the future** (e.g. weather forecast): $y =$ temperature or amount of rain


Regression in Machine Learning: Applications

Domains of application:

- ▶ Price estimation/prediction
- ▶ Weather forecast
- ▶ Production quantity estimation
- ▶ Stock option price prediction
- ▶ Fit statistical model to data
- ▶ Physics & chemistry
- ▶ ... and others

Linear and polynomial regression

Purpose of regression: **approximate solutions** of **overdetermined systems**.



In this course, we will see

- ▶ Linear regression
- ▶ Polynomial regression

Linear regression

Linear regression

Principal components:

- ▶ Old problem (least-squares method usually credited to Carl Friedrich Gauss in 1795)

Linear regression

Principal components:

- ▶ Old problem (least-squares method usually credited to Carl Friedrich Gauss in 1795)
- ▶ Several ways to approximate the data
 - ▶ Linear model
 - ▶ Polynomial model (remember kernels from SVMs)
 - ▶ Fit a distribution
 - ▶ ...

Linear regression

Principal components:

- ▶ Old problem (least-squares method usually credited to Carl Friedrich Gauss in 1795)
- ▶ Several ways to approximate the data
 - ▶ Linear model
 - ▶ Polynomial model (remember kernels from SVMs)
 - ▶ Fit a distribution
 - ▶ ...
- ▶ Several ways to formulate the problem
 - ▶ Least Squares
 - ▶ Support Vector regression
 - ▶ ...

Linear regression

Principal components:

- ▶ Old problem (least-squares method usually credited to Carl Friedrich Gauss in 1795)
- ▶ Several ways to approximate the data
 - ▶ Linear model
 - ▶ Polynomial model (remember kernels from SVMs)
 - ▶ Fit a distribution
 - ▶ ...
- ▶ Several ways to formulate the problem
 - ▶ Least Squares
 - ▶ Support Vector regression
 - ▶ ...
- ▶ Several ways to solve the problem
 - ▶ Closed-form expression (exact formula)
 - ▶ Optimization

Linear regression: Toy example

living area (m ²)	# bedrooms	price (1000's euros)
50	1	30
76	2	48
26	1	12
102	3	90

Linear regression: Toy example

living area (m ²)	# bedrooms	price (1000's euros)
50	1	30
76	2	48
26	1	12
102	3	90
61	2	?

Linear regression: Toy example

living area (m ²)	# bedrooms	price (1000's euros)
50	1	30
76	2	48
26	1	12
102	3	90
61	2	?

Linear model: $\text{price} = \mathbf{a} \times \text{area} + \mathbf{b} \times \# \text{ bedrooms} + \mathbf{c}$

Linear regression: Toy example

living area (m ²)	# bedrooms	price (1000's euros)
50	1	30
76	2	48
26	1	12
102	3	90
61	2	?

Linear model: $\text{price} = \mathbf{a} \times \text{area} + \mathbf{b} \times \# \text{ bedrooms} + \mathbf{c}$

Problem: optimal values for **a**, **b** and **c**?

Linear regression: Toy example

living area (m ²)	# bedrooms	price (1000's euros)
50	1	30
76	2	48
26	1	12
102	3	90
61	2	?

Linear model: $\text{price} = \mathbf{a} \times \text{area} + \mathbf{b} \times \# \text{ bedrooms} + \mathbf{c}$

Problem: optimal values for **a**, **b** and **c**?

General formulation:

$$\hat{y} = w^T x$$

Linear regression with ordinary least-squares

Linear regression: Estimate y as a **linear** function of x :

$$\hat{y} = w^T x$$

Linear regression with ordinary least-squares

Linear regression: Estimate y as a **linear** function of x :

$$\hat{y} = w^T x$$

Least squares: Penalty (loss) is a **quadratic** function

$$\ell(\hat{y}, y) = (\hat{y} - y)^2$$

Regression formulation

2 main ways to solve the linear least-squares problem:

$$\min_w \sum_i \left(w^T x^{(i)} - y^{(i)} \right)^2 \quad (1)$$

Method 1: Closed-form expression

$$w = \left(X^T X \right)^{-1} X^T Y$$

It can be computationally expensive (matrix multiplication, matrix inversion, matrix multiplication, matrix-vector multiplication)

Method 2: Numerical optimization For example gradient descent algorithms, ...

Polynomial regression

Polynomial regression \approx Kernel trick

Remind the kernel trick from the SVM lecture.

Example:

$$x = (x_1, x_2)$$

can become

$$\phi(x) = (1, x_1, x_2, x_1^2, x_2^2, x_1x_2)$$

with a **second order kernel**.

Then we can find w solution of

$$\min_w \sum_i \left(w^T \phi \left(x^{(i)} \right) - y^{(i)} \right)^2 \quad (2)$$

Practical information

Variable standardization

Variables have various magnitudes. Example:

- ▶ Living area: Up to a few hundreds m^2
- ▶ Price: Up to a few 100 000s BGN (and even more)
- ▶ # bedrooms: usually much smaller than 10

This can be an issue when training a regression model.

Variable standardization

Variables have various magnitudes. Example:

- ▶ Living area: Up to a few hundreds m²
- ▶ Price: Up to a few 100 000s BGN (and even more)
- ▶ # bedrooms: usually much smaller than 10

This can be an issue when training a regression model.

It is possible to calculate the **standard score** z of a variable x

$$z = \frac{x - \mu}{\sigma}$$

where

- ▶ μ is the mean of the variable
- ▶ σ is its standard deviation

Variable standardization

Variables have various magnitudes. Example:

- ▶ Living area: Up to a few hundreds m²
- ▶ Price: Up to a few 100 000s BGN (and even more)
- ▶ # bedrooms: usually much smaller than 10

This can be an issue when training a regression model.

It is possible to calculate the **standard score** z of a variable x

$$z = \frac{x - \mu}{\sigma}$$

where

- ▶ μ is the mean of the variable
- ▶ σ is its standard deviation

Another option: Scale between 0 and 1

$$z = \frac{x - \min}{\max - \min}$$

Overfitting and underfitting

Illustration on a generated example: Try to fit the function

$$y = f(x) = \cos\left(\frac{3\pi}{2}x\right) + \text{noise}$$




for $x \in [0, 1]$, with a polynomial regression

Overfitting and underfitting

Illustration on a generated example: Try to fit the function

$$y = f(x) = \cos\left(\frac{3\pi}{2}x\right) + \text{noise}$$

for $x \in [0, 1]$, with a polynomial regression



Parameter selection

As for classification models, parameter selection plays a key role in the regression performance:

- ▶ Degree of the polynomial
- ▶ Regularization parameter

Parameter selection

As for classification models, parameter selection plays a key role in the regression performance:

- ▶ Degree of the polynomial
- ▶ Regularization parameter

Optimal parameters can be chosen with cross-validation over a grid:

Parameter selection

As for classification models, parameter selection plays a key role in the regression performance:

- ▶ Degree of the polynomial
- ▶ Regularization parameter

Optimal parameters can be chosen with cross-validation over a grid:

- ▶ Split the data into train/test

Parameter selection

As for classification models, parameter selection plays a key role in the regression performance:

- ▶ Degree of the polynomial
- ▶ Regularization parameter

Optimal parameters can be chosen with cross-validation over a grid:

- ▶ Split the data into train/test
- ▶ Choose a degree $d \in \{1, \dots, 20\}$

Parameter selection

As for classification models, parameter selection plays a key role in the regression performance:

- ▶ Degree of the polynomial
- ▶ Regularization parameter

Optimal parameters can be chosen with cross-validation over a grid:

- ▶ Split the data into train/test
- ▶ Choose a degree $d \in \{1, \dots, 20\}$
- ▶ Train on the train set with this degree

Parameter selection

As for classification models, parameter selection plays a key role in the regression performance:

- ▶ Degree of the polynomial
- ▶ Regularization parameter

Optimal parameters can be chosen with cross-validation over a grid:

- ▶ Split the data into train/test
- ▶ Choose a degree $d \in \{1, \dots, 20\}$
- ▶ Train on the train set with this degree
- ▶ Test the model on the test set

Parameter selection

As for classification models, parameter selection plays a key role in the regression performance:

- ▶ Degree of the polynomial
- ▶ Regularization parameter

Optimal parameters can be chosen with cross-validation over a grid:

- ▶ Split the data into train/test
- ▶ Choose a degree $d \in \{1, \dots, 20\}$
- ▶ Train on the train set with this degree
- ▶ Test the model on the test set


It can be done over several train/test splits.

Toy example: Fitting a distribution

Find A , x_0 and σ such that

$$\hat{y} = f(x) = Ae^{-\frac{(x - x_0)^2}{2\sigma^2}}$$

best fits the data in terms of least-square error.



Regression with SVMs

Alternatives to least squares

It is possible to use a different loss function ℓ . Remember, we had

$$\ell(\hat{y}, y) = (\hat{y} - y)^2$$

Alternatives to least squares

It is possible to use a different loss function ℓ . Remember, we had

$$\ell(\hat{y}, y) = (\hat{y} - y)^2$$

We can use **support vector machines for regression** (SVR):

- ▶ If **within the margin** (i.e. $-\epsilon \leq \hat{y} - y \leq +\epsilon$) then **no penalty**
- ▶ linear or quadratic **penalty outside the margin**

(see flip-chart for illustration)

This loss function is called ϵ -insensitive.

Alternatives to least squares

It is possible to use a different loss function ℓ . Remember, we had

$$\ell(\hat{y}, y) = (\hat{y} - y)^2$$

We can use **support vector machines for regression** (SVR):

- ▶ If **within the margin** (i.e. $-\epsilon \leq \hat{y} - y \leq +\epsilon$) then **no penalty**
- ▶ linear or quadratic **penalty outside the margin**

(see flip-chart for illustration)

This loss function is called ϵ -insensitive.

Note: We can use kernels as for SVM

Conclusion

Regression: Output y is a **continuous variable**.

Conclusion

Regression: Output y is a **continuous variable**.

Several ways to **penalize errors**:

- ▶ Least squares
- ▶ Support vector regressions

Conclusion

Regression: Output y is a **continuous variable**.

Several ways to **penalize errors**:

- ▶ Least squares
- ▶ Support vector regressions

Several ways to **model the prediction**:

- ▶ Linear
- ▶ Quadratic
- ▶ Other kernel

Conclusion

Regression: Output y is a **continuous variable**.

Several ways to **penalize errors**:

- ▶ Least squares
- ▶ Support vector regressions

Several ways to **model the prediction**:

- ▶ Linear
- ▶ Quadratic
- ▶ Other kernel

Parameter selection is important

Thank you! Questions?